

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation



# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.00	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.00 to 10.15		Break
10.15 to 11.00	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.00 to 11.20		Break
11.20 to 12.15	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.15 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.00	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.00 to 10.15		Break
10.15 to 11.00	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.00 to 11.20		Break
11.20 to 12.15	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.15 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# Tools from Intel

- Speaker: Harish Patil
  - Principal Engineer, Intel Corporation
- Topics Covered
  - Binary instrumentation using Pin or writing Pintools
  - PinPlay kit and PinPlay-enabled tools
  - SDE build kit for microarchitecture emulation
  - Checkpointing threaded applications using PinPlay, SDE
  - Detailed discussion on ELFies including its generation and usage

# Simulation and Sampling Overview

- Speaker: Akanksha Chaudhari
  - Research Assistant, National University of Singapore
- Topics Covered
  - Architectural exploration and evaluation
  - Simulation as a tool for performance estimation
  - Methods for fast estimation using simulation
  - State-of-the-art single-threaded sampled simulation techniques

# LoopPoint Methodology

- Speaker: Alen Sabu
  - PhD Candidate, National University of Singapore
- Topics Covered
  - Sampled simulation of multi-threaded applications
  - Existing methodologies and their drawbacks
  - Detailed discussion on LoopPoint methodology
  - Experimental results of LoopPoint

# Simulation and Demo

- Speaker: Changxi Liu
  - PhD Student, National University of Singapore
- Topics Covered
  - Overview of Sniper simulator
  - High-level structure of LoopPoint code
  - Demo on how to use LoopPoint tools
  - Integrating workloads to run with LoopPoint

# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.00	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.00 to 10.15	Break	
10.15 to 11.00	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.00 to 11.20	Break	
11.20 to 12.15	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.15 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation

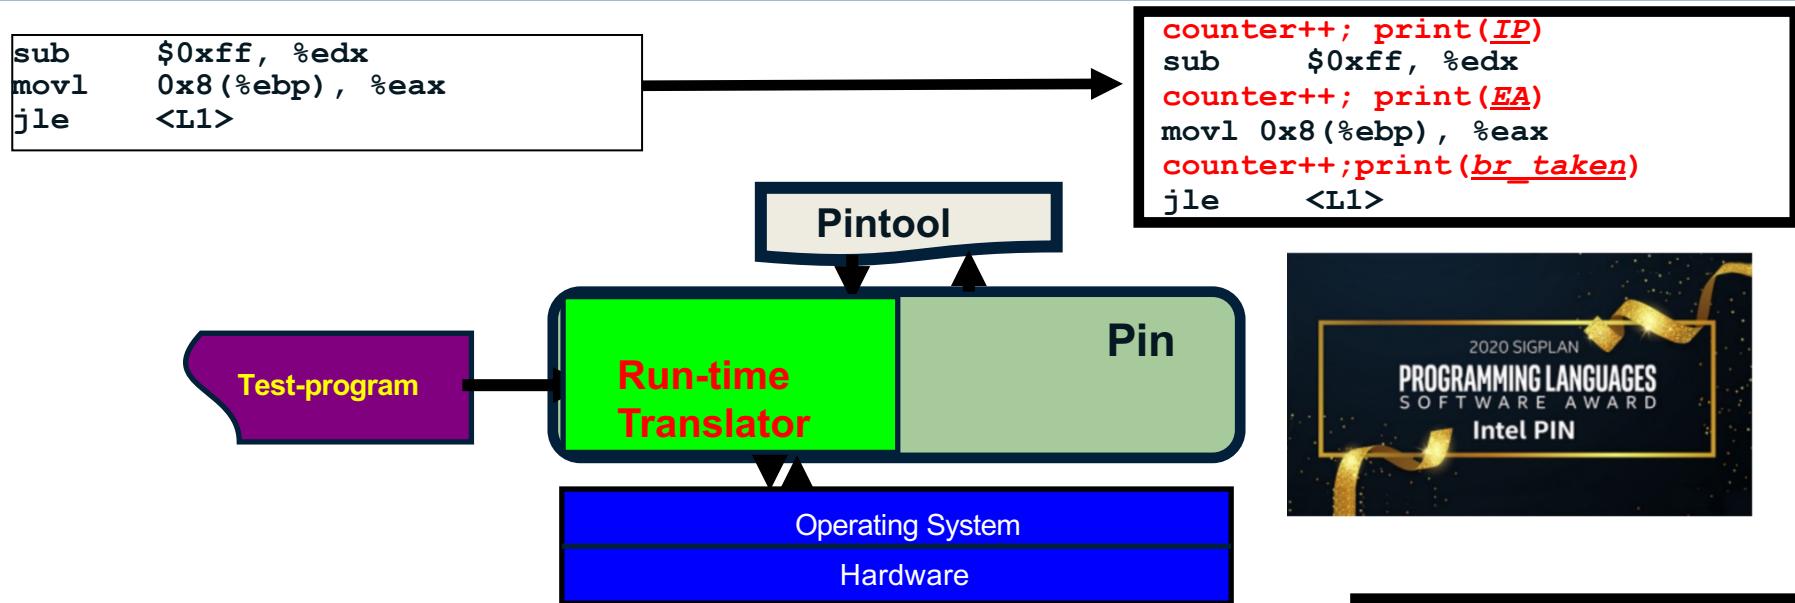


Session 1

# Tools and Methodologies

HARISH PATIL, PRINCIPAL ENGINEER (DEVELOPMENT TOOLS SOFTWARE)  
INTEL CORPORATION

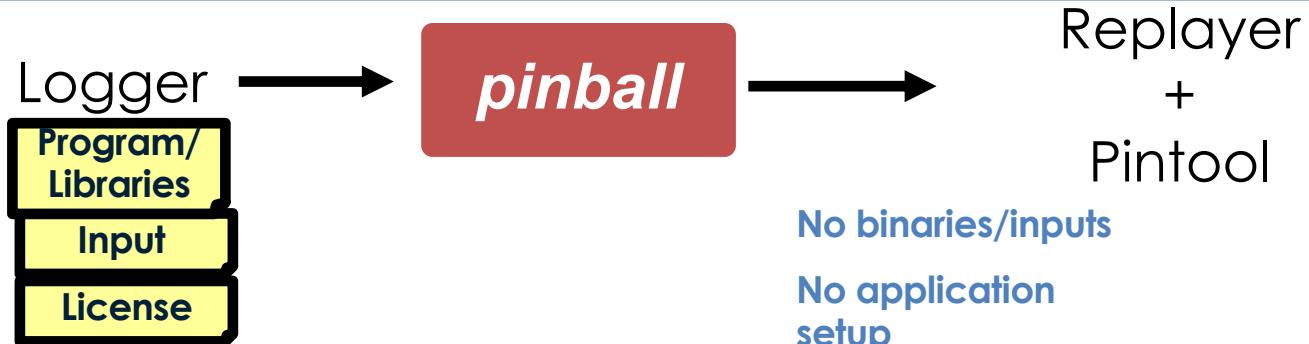
# Pin: A Tool for Writing Program Analysis Tools



Normal output +  
*Analysis output*



# *PinPlay*: Software-based User-level Capture and Replay



**Platforms** : Linux, Windows, MacOS

No binaries/inputs


**Upside** : It works! Large OpenMP / MPI programs, Oracle

No application setup

**Downside** : High run-time overhead: ~100-200X for capture →  
Cannot be turned on all the time

No license  
checking

# *Pinball (single-threaded):* Initial memory/register + injections



foo.text

- **System calls** : skipped by injecting next rip/ memory changed
- **CPUID, RDTSC** : affected registers injected
- **Signals/Callbacks** : New register state injected

# *Pinball (multi-threaded):* Pinball (single-threaded) + Thread-dependencies

foo.reg (per-thread)



foo.text

Application Memory (common)

foo.reg (per-thread)

Event injection works only if same behavior  
(same instruction counts) is guaranteed  
during replay

foo.sel (per-thread)

[T1] 2 T2 2  
[T1] 3 T2 3

[T2] 5 T4 1


Thread T2 cannot execute instruction  
5 until T4 executes instruction 1

foo.race (per-thread)

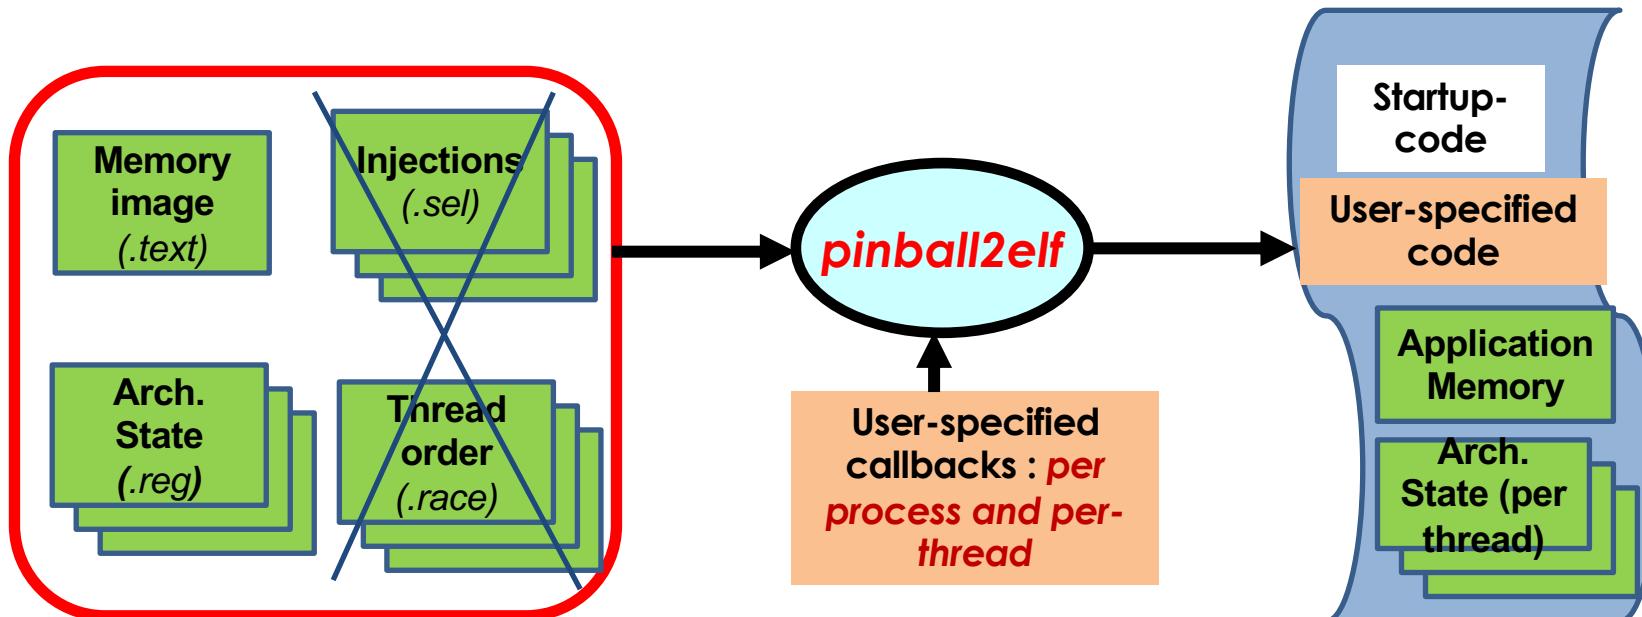
Thread T1 cannot execute instruction 2  
until T2 executes instruction 2

# *ELFie : An Executable Application Checkpoint*

- **Checkpoint**: Memory + Registers
- **Application** : Only program state captured -- no OS or simulator states
- **Executable** : In the Executable Linkage Format commonly used on Linux



# *pinball2elf*: Pinball converter to ELF



# Getting started with *pinball2elf*

**Prerequisite:** ‘*perf*’ installed on your Linux box (*perf stat /bin/ls* should work)

- Clone pinball2elf repository: *git clone https://github.com/intel/pinball2elf.git*
- *cd pinball2elf/src*
- *make all*
- *cd .../examples/ST*
- *./testST.sh*

Running ..../scripts//pinball2elf.basic.sh pinball.st/log\_0

..

Running ..../scripts//pinball2elf.perf.sh pinball.st/log\_0 st  
export ELFIE\_PERFLIST=0:0,0:1,1:1

...

hw\_cpu\_cycles:47272 hw\_instructions:4951 sw\_task\_clock:224943

---

*Tested : Ubuntu 20.04.4 LTS : gcc/g++ 7.5.0 and 9.4.0  
and Ubuntu 18.04.6 LTS: gcc/g++ 7.5.0*

# *ELFie types: basic, sim, perf*

	<i>basic</i>	<i>sim</i>	<i>perf</i>
How to create	<code>scripts/pinball2elf.basic.sh pinball</code>	<code>scripts/pinball2elf.sim.sh pinball</code>	<code>scripts/pinball2elf.perf.sh pinball perf.out</code>
Exits gracefully?	NO, either hangs or dumps core	NO, either hangs or dumps core Simulator handles exit	YES, when retired instruction count reaches pinball icount
Environment variables used	NONE	<code>ELFIE_VERBOSE=0/1</code> <code>ELFIE_COREBASE=X</code> Set affinity : thread 0 → core X, thread 1 → core x+1	" <code>ELFIE_WARMUP</code> " to decide whether to use warmup " <code>ELFIE_PCCONT</code> " to decide how to end warmup/simulation regions <code>ELFIE_PERFLIST</code> , enables performance counting

# Example: *ELFIE\_PERFLIST* with a *perf ELFie*

*ELFIE\_PERFLIST*, enables performance counting

( based on /usr/include/linux/perf\_event.h  
perf type: 0 --> HW 1 --> SW  
HW counter: 0 --> PERF\_COUNT\_HW\_CPU\_CYCLES  
HW counter: 1 --> PERF\_COUNT\_HW\_CPU\_INSTRUCTIONS  
SW counter: 0 --> PERF\_COUNT\_SW\_CPU\_CLOCK  
... <see perf\_event.h:'enum perf\_hw\_ids' and 'enum

perf\_sw\_ids')

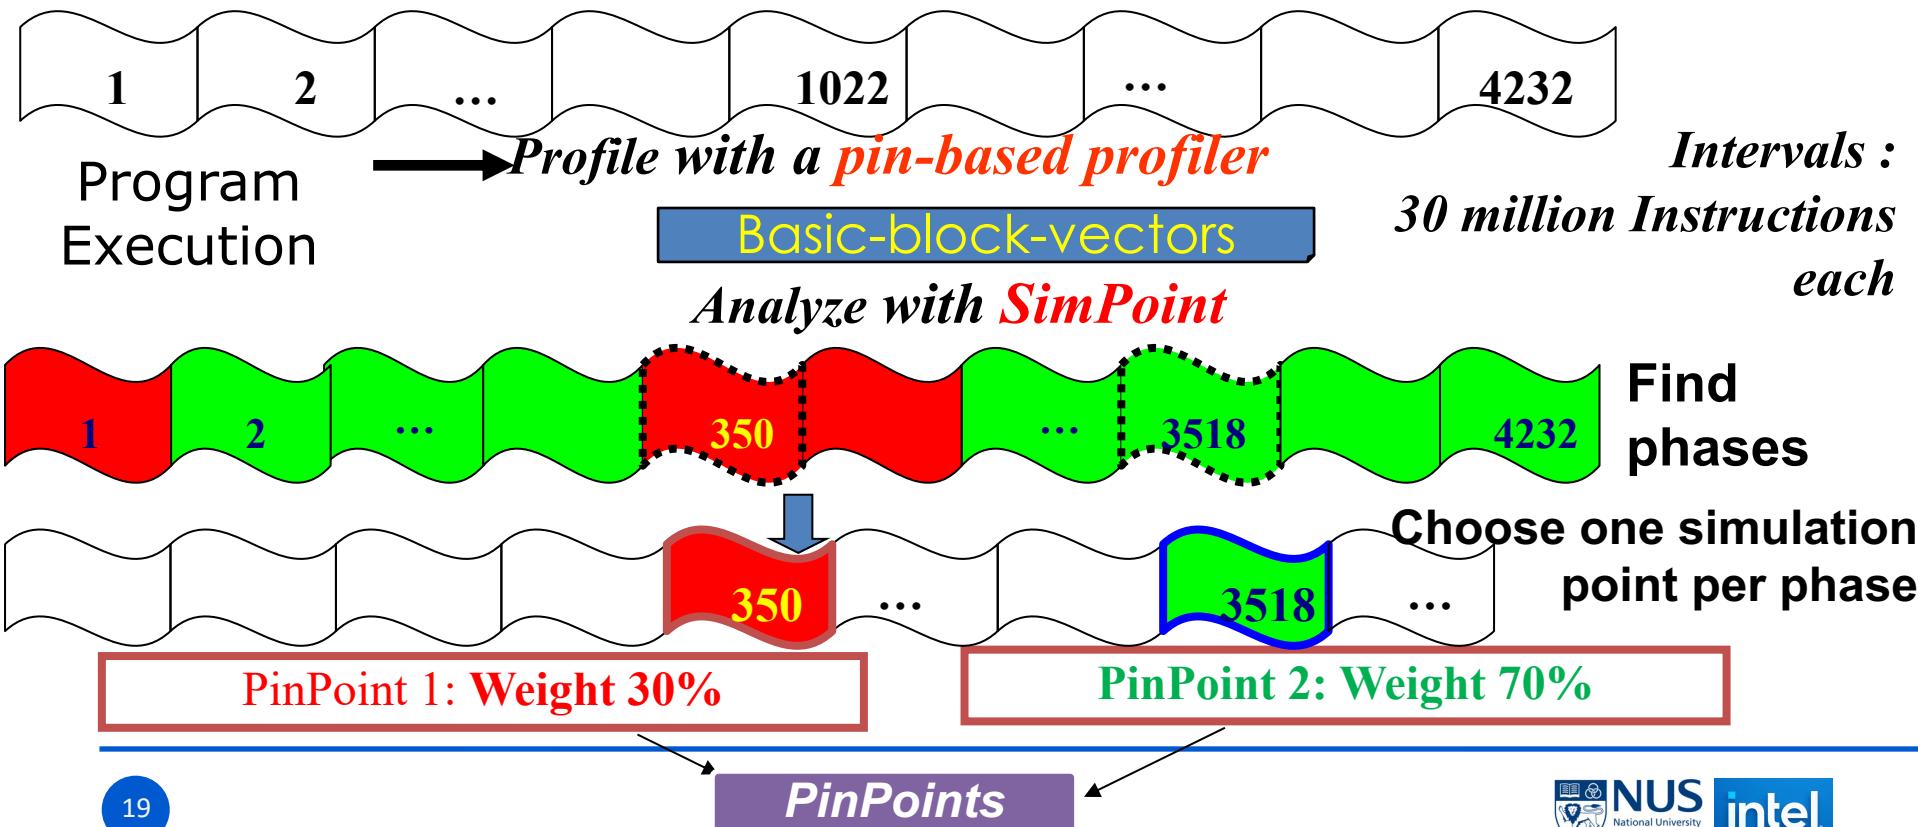
```
% cd examples/MT  
% ../../scripts/pinball2elf.perf.sh pinball.mt/log_0 perf.out  
% setenv ELFIE_PERFLIST "0:0,0:1,1:1"
```

```
% pinball.mt/log_0.perf.elfie
```


```
└── perf.out.0.perf.txt  
└── perf.out.1.perf.txt  
└── perf.out.2.perf.txt
```

```
ROI start: TSC 48051110586217756  
Thread start: TSC 48051110623843452  
-----  
Simulation end: TSC 48051110625045322  
Sim-end-i-count 3436  
hw_cpu_cycles:36148 hw_instructions:3476  
sw_task_clock:141901  
-----  
Thread end: TSC 48051110625366502  
ROI end: TSC 48051110625959364  
hw_cpu_cycles:40097 hw_instructions:4455  
sw_task_clock:188637
```

# *PinPoints* == *Pin + SimPoint*



# *PinPoints* : The repeatability challenge

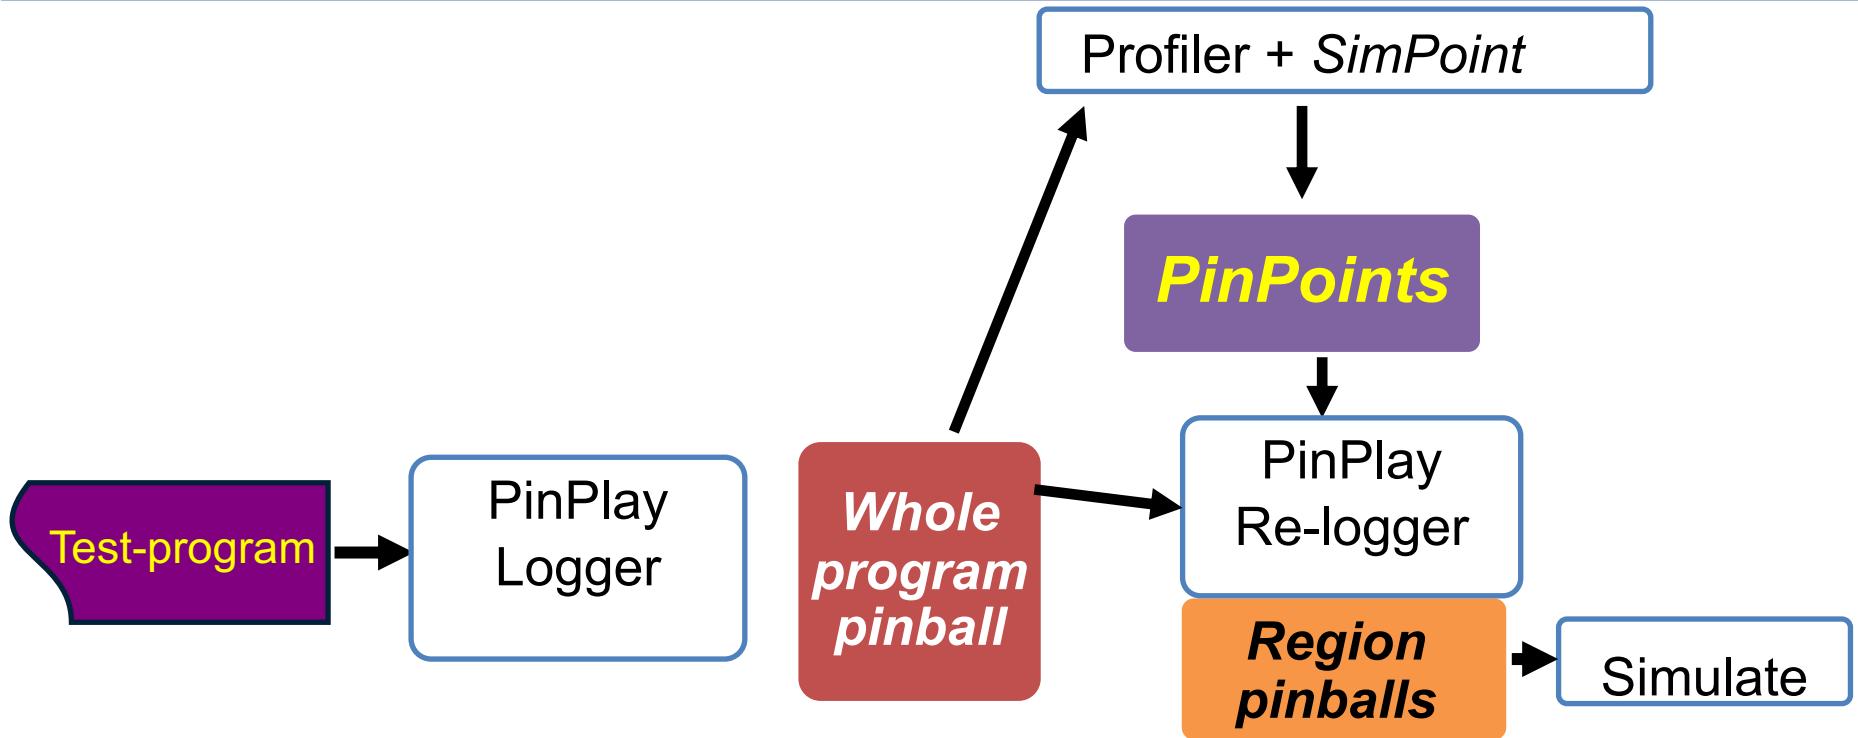


**Problem:** Two runs are not exactly same → PinPoints missed (PC marker based)

[ “*PinPoints out of order*” “*PinPoint End seen before Start*” ]

Found this for 25/54 SPEC2006 runs!

# *PinPlay* provides repeatability



# Single-threaded *PinPoints* → SPEC2006/2017 pinballs publicly available

1. University of California (San Diego), Intel Corporation, and Ghent University

<https://www.spec.org/cpu2006/research/simpoint.html>

2. University of Texas at Austin

<https://www.spec.org/cpu2017/research/simpoint.html>

3. Northwestern University

[Public Release and Validation of SPEC CPU2017 PinPoints](#)

# *DCFG* Generation with *PinPlay*

## Dynamic Control-Flow Graph (DCFG)


Directed graph extracted for a specific execution:

Nodes → basic blocks


Edges → control-flow : augmented with per-thread execution counts



# *PinPlay + DCFG : Stronger Repeatability*



# *LoopPoint*: Simulation alternatives




**Requirement:** Execution invariant region specification  
(PC+count for compute loop entries)

# Intel Software Development Emulator (*Intel SDE*)

- The Intel® Software Development Emulator is a **functional user-level (ring 3) emulator** for x86 (32b and 64b) new instructions built upon Pin and XED (X86 encoder/decoder)
- **Goal:** New instruction/register emulation between the time when they are designed and when the hardware is available.
- Used for compiler development, architecture and workload analysis, and tracing for architecture simulators
- No special compilation required
- Supported on Windows/Linux/Mac OS
- Runs only in user space (ring 3)

# How SDE Works

- Based on Pin (<http://pintool.intel.com>) and XED decoder/encoder (<https://github.com/intelxed/xed>)
- Instrument new instructions
  - Add call to emulation routine
  - Delete original instruction
- Emulation routine:
  - Update native state with emulated state



# Using *SDE* for *PinPoints* and *LoopPoint*

Prerequisites:

1. SDE build kit (version 9.0 or higher) from Intel

<http://www.intel.com/software/sde>

2. pinplay-tools from Intel

<https://github.com/intel/pinplay-tools>

3. SimPoint sources from UCSD

<https://cseweb.ucsd.edu/~calder/simpoint/>

4. Pinball2elf sources from Intel

<http://pinelfie.org> → <https://github.com/intel/pinball2elf>

# Getting ready for *LoopPoint* ...

1. Expand SDE build-kit : *setenv SDE\_BUILD\_KIT<path to SDE kit>*
2. *cp -r pinplay-tools/pinplay-scripts \$ SDE\_BUILD\_KIT*
3. Build simpoint (see pinplay-tools/pinplay-scripts/README.simpoint)
  - *cp <path>/SimPoint.3.2/bin/simpoint \$ SDE\_BUILD\_KIT/pinplay-scripts/PinPointsHome/Linux/bin/*
4. Build global looppoint tools
  - *setenv PINBALL2ELF <path to pinball2lef repo>*
  - *cd pinplay-tools/GlobalLoopPoint*
  - *./sde-build-GlobalLoopPoint.sh*

# **SDE kit expanded for LoopPoint**

**sde-external-9.0.0-2021-11-07-lin**



# Running *LoopPoint* for an OpenMP program

- `cd pinplay-tools/dotproduct-omp` # see README there
- `make` # builds dotproduct-omp → base.exe
- `./sde-run.looppoint.global_looppoint.concat.filter.flowcontrol.sh`

`~/pinplay-tools/dotproduct-omp`

  └── `dotproduct.1_282016.Data`

  └── `dotproduct.1_282016.pp`

  └── `whole_program.1`

bbv files (\*.bb), PinPoints  
file (\*.csv, \*.CSV)

Region pinballs

Whole-program pinball + DCFG

# Summary: Simulation of Multi-threaded Programs: Tools & Methodologies

Where to simulate?

**SDE + LoopPoint**  
Compute-loop iterations as  
'Unit of work'

How to simulate?

1. Pinball-driven
2. ELFie-driven
3. Binary-driven

Are the regions representative?

1. Simulation (Sniper) -based  
2. ELFie-based / Binary+ROI/Perf (*not covered*)  
Whole-program performance vs  
Region-predicted performance

# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.30	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.30 to 10.45		Break
10.45 to 11.30	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.30 to 11.40		Break
11.40 to 12.20	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.20 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation



Session 2

# Simulation and Sampling

AKANKSHA CHAUDHARI, RESEARCH ASSISTANT  
NATIONAL UNIVERSITY OF SINGAPORE

# Architectural Trends in Processor Design





Fig. 1: Moore Law number of transistor per device: past, present, future  
[Intel]

- Moore's Law predicts that the number of transistors per device will double every two years.
- First microprocessor had 2200 transistors – Intel aspiring to have 1 trillion transistors by 2030.

# Architectural Trends in Processor Design




Main Goal: Meeting the ever-increasing computational demands *while* adhering to stringent non-functional requirements (ex: size, power)!

Fig. 2: Transistor innovations over time

# Exploration and Evaluation of New Ideas



- Architecture is rapidly evolving domain with a lot of new research directions.
- A plethora of design choices are available:
  - Ranging from the choice of components, the choice of operating modes of each component, the choice of interconnects used, the choice of algorithms employed, etc.
- The process of exploration and evaluation of new ideas is often complex and time-consuming.

# Exploration and Evaluation of New Ideas



# Exploration and Evaluation of New Ideas

The Architect IRL



The Important Question:

So how do we then **explore new ideas quickly** and evaluate them accurately to find the **BEST** idea?

# Exploration and Evaluation of New Ideas

- Important questions when considering any idea:
  - Does it work?
  - How well does it work?
- Generally speaking, good idea optimizes a finite set of performance metrics, say M.

Let  $M = \{m_1, m_2, \dots, m_n\}$ ,

where  $m_i \in M$  can be computational speed, energy efficiency,  
memory utilization etc.

# Exploration and Evaluation of New Ideas

A variety of different evaluation methods are available:

- Theoretical proofs
  - Formally proving the correctness/efficiency of the proposed design using computational theory and mathematical logic.
- Analytical modeling
  - Mathematically modeling the proposed design at some level of abstraction to analyze/quantify its performance.

# Exploration and Evaluation of New Ideas

A variety of different evaluation methods are available:

- Simulation (at varying degrees of abstraction and accuracy)  
A model that mimics the system behavior demonstrating its key functions and operations accurately.
- Prototyping using existing systems  
Implementing a draft version of the proposed design using existing systems (such as FPGAs) to evaluate the performance.
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Theoretical proofs / accurate modeling of practical systems can be extremely complex

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Theoretical proofs / accurate modeling of practical systems can be extremely complex

Difficult to evaluate a design for all possible workload profiles

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Theoretical proofs / accurate modeling of practical systems can be extremely complex

Difficult to evaluate a design for all possible workload profiles

Worst-case estimates can be misleading

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Note that, using analytical modeling in conjunction with simulation can provide significant quantifiable benefits

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Relatively Expensive

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Relatively Expensive

Limited by the capability  
of the systems used to  
model

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Most expensive

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

Most expensive

Not always feasible to implement  
and evaluate each idea  
Especially if we have too many  
options to choose from

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

# Exploration and Evaluation of New Ideas

An “evaluation” of the evaluation methods:

- Theoretical proof
- Analytical modelling
- Simulation
- Prototyping
- Actual implementation

The most feasible way to explore and evaluate large-scale, complex architectural designs in terms of time, cost and efficiency!

# Simulation: An Overview

- Simulation enables the modeling of new research ideas at varying degrees of abstraction and accuracy.
- Main goals:
  - Enables fast exploration of design space (to discover the next big idea!).
  - Evaluation of new research ideas by estimating their relative performances.
  - Evaluation, debugging and understanding the behavior of existing systems.
- How does a simulator work?
  - Mimics system behavior to reflect its performance in terms of the metric of interest (ex: Instructions per cycle, Runtime, etc).


# Simulation: An Overview

- Caution: Inaccurate simulation → Inaccurate evaluation/results → Wrong conclusions.
  - Ex: Inaccurate assumptions, inaccurate extrapolation of performance, etc.
- Very important to select the right simulation technique!

# Simulation: An Overview

- An ideal simulation technique:
  - High speed: For faster exploration.
  - High flexibility: For wider exploration.
  - High accuracy/low simulation error: For accurate evaluation.
- Practical simulation techniques → tradeoffs:
  - Speed vs. accuracy
  - Speed vs. flexibility
  - Flexibility vs. accuracy

# Different Simulation Techniques



# Techniques to Simulate Faster

- Partially simulating to extrapolate performance:

- Simulating the first 1 billion instructions in detail.



- Fast-forwarding to skip the initialization phase and then simulating 1 billion instructions in detail.



- Fast-forwarding to skip the initialization phase, microarchitectural state warming, and then simulating the 1 billion instructions in detail



# Techniques to Simulate Faster


- Workload reduction
  - Simulating for reduced input sets
  - Simulating for reduced loop counts in workloads

# Techniques to Simulate Faster

- Workload reduction
  - Simulating for reduced input sets
  - Simulating for reduced loop counts in workloads
- Problems with these techniques:

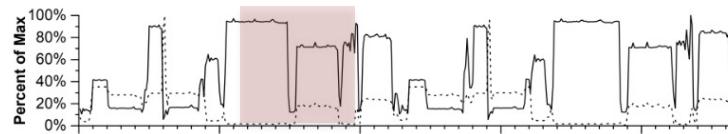
# Techniques to Simulate Faster

- Workload reduction
  - Simulating for reduced input sets
  - Simulating for reduced loop counts in workloads
- Problems with these techniques:
  - [Partial simulation + extrapolation] → fail to capture global variations in program behavior and performance.




# Techniques to Simulate Faster

- Workload reduction
  - Simulating for reduced input sets
  - Simulating for reduced loop counts in workloads
- Problems with these techniques:
  - [Partial simulation + extrapolation] → fail to capture global variations in program behavior and performance.



- [Workload reduction] → benchmark behavior varies significantly across test, train and reference inputs → do not reflect the actual performance.

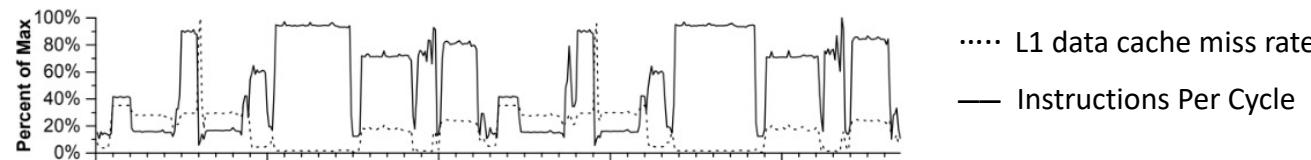
# Sampled Simulation to the Rescue!

- Sampling enables the simulation selective representative regions of an application.
  - “Representative regions” refer to the subset of regions in the application that reflect the behavior of the entire system when extrapolated.
- How to select these “representative regions”?
  - Targeted sampling (like in SimPoint)  

    - (Full) program execution
    - Representative regions
  - Statistical sampling (like in SMARTS)  

    - Representative regions

# Sampled Simulation Techniques: SimPoint

- Large-scale program behaviors vary significantly over their run times.
  - Difficult to estimate performance using previously discussed techniques.



- Main idea behind SimPoint:
  - Automatically & efficiently analyzing program behavior over different phases of execution.
- SimPoint uses Basic Block Vectors (BBV) as a hardware-independent metric for characterizing the program behavior in different phases.

# Sampled Simulation Techniques: SimPoint


- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions

# Sampled Simulation Techniques: SimPoint

- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions

# Sampled Simulation Techniques: SimPoint

A Basic Block Vector (BBV) is a single-dimensional array that maintains a count of how many times each basic block was run in a given interval during the program execution.



# Sampled Simulation Techniques: SimPoint

- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions


# Sampled Simulation Techniques: SimPoint

- Basic Block Similarity: Measured using Euclidean or Manhattan Distances.



$$EuclideanDist(a, b) = \sqrt{\sum_{i=1}^D (a_i - b_i)^2}$$

$$ManhattanDist(a, b) = \sum_{i=1}^D |a_i - b_i|$$

- Depicted by Basic Block Similarity Matrices.



Using Manhattan distances



Using Euclidian distances


- The program execution progresses along the diagonal of the matrix.
- Point at  $(x, y)$  gives similarity index between  $BBV_x$  and  $BBV_y$ .
- ↑ darkness → ↑ similarity

# Sampled Simulation Techniques: SimPoint

- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions

# Sampled Simulation Techniques: SimPoint

- The Basic Block Vectors obtained from the basic block profiling step have a very large number of dimensions! (in the range of 2,000 -- 100,000)
- “Curse of dimensionality”:
  - Hard to cluster data as the number of dimensions increases.
  - Clustering time increases significantly wrt as the number of dimensions increases.
- Solution: Reduce the number of dimensions to 15 using Random Linear Projections.



# Sampled Simulation Techniques: SimPoint


- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions

# Sampled Simulation Techniques: SimPoint

K-means clustering:

- Initialize k cluster centers by randomly choosing k points from the data.
- Repeat until convergence:
  - Do for all data points:
    - Compare the distance from all k cluster centers.
    - Assign it to the cluster with the closest center.
  - Update cluster center to the centroid of the newly assigned memberships.

*Choosing k: The clustering that achieves a BIC score that is at least 90% of the spread between the largest and smallest BIC score is chosen.*




# Sampled Simulation Techniques: SimPoint

- How SimPoint works:
  - STEP 1: Basic block profiling
    - Generating the Basic Block Vectors
    - Creating a Basic Block Similarity Matrix
  - STEP 2: Clustering of Basic Block Vectors
    - Random Projection
    - K-means Clustering
  - STEP 3: Identifying representative regions

# Sampled Simulation Techniques: SimPoint

- Representative region → single simulation point
  - BBV with the lowest distance from the centroid of all cluster centers.
- Representative regions → multiple simulation points
  - For each cluster, choose the BBV that is closest to the centroid of the cluster.



# Sampled Simulation Techniques: SMARTS

- Main idea behind SMARTS:
  - Using systematic sampling:
    - To identify a minimal but representative sample from the population for microarchitecture simulation
    - To establish a confidence level for the error on sample estimates
  - Simulating using two modes :
    - Detailed simulation of sampled instructions → accounting for all the microarchitectural details.
    - Functional simulation of remaining instructions → accounting only for the programmer-visible architectural states (ex: registers, memory).

# Sampled Simulation Techniques: SMARTS

- STEP 1: Determine n based on the required confidence (assuming the coefficient of variation  $V_x$ ) using the following equation:

$$\text{confidence interval} = \mp \left[ \frac{z \cdot V_x}{\sqrt{n}} \right] \cdot X$$

(where  $X$  is the mean, and  $z = 100 (1 - \alpha/2)$  is the percentile of the standard normal distribution)


- STEP 2: If the initial sample does not achieve the desired confidence, compute n using the equation:

$$n \geq \left[ \left( \frac{z \cdot V_x'}{\varepsilon} \right)^2 \right]$$

(where  $V_x'$  is the coefficient of variation for the obtained sample)


# Sampled Simulation Techniques: SMARTS

- SMARTS uses Systematic Sampling:



# Sampled Simulation Techniques: SMARTS

- Simulation:



# Sampled Simulation Techniques: SMARTS

- Evaluation results:
    - Average error:
      - 0.64% for CPI
      - 0.59% for EPI
    - Speedup over full-stream simulation:
      - 35x for 8-way out-of-order processors
      - 60x for 16-way out-of-order processors
- } By simulating fewer than 50 million instructions in detail per benchmark.

# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.30	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.30 to 10.45	Break	
10.45 to 11.30	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.30 to 11.40	Break	
11.40 to 12.20	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.20 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation



Session 3


# The LoopPoint Methodology

ALEN SABU, PHD CANDIDATE

NATIONAL UNIVERSITY OF SINGAPORE

# Simulation in the Post-Dennard Era


- Modern architectures require smarter simulators
- Microarchitectural simulation is slow
  - NPB (D), SPEC CPU2017 (ref) can take years
  - Solution – Simulate representative sample



Intel's Alder Lake die shot.  
Image source: WikiChip

# Simulation in the Post-Dennard Era

- Modern architectures require smarter simulators
- Microarchitectural simulation is slow
  - NPB (D), SPEC CPU2017 (ref) can take years
  - Solution – Simulate representative sample




Benchmarks with 8 threads, static schedule,  
passive wait-policy, simulated at 100 KIPS.

# Simulation in the Post-Dennard Era

- Modern architectures require smarter simulators
- Microarchitectural simulation is slow
  - NPB (D), SPEC CPU2017 (ref) can take years
  - Solution – Simulate representative sample



Can we further bring  
down simulation time




Benchmarks with 8 threads, static schedule,  
passive wait-policy, simulated at 100 KIPS.

# Simulation in the Post-Dennard Era

- Modern architectures require smarter simulators
- Microarchitectural simulation is slow
  - NPB (D), SPEC CPU2017 (ref) can take years
  - Solution – Simulate representative sample



Can we further bring  
down simulation time



Benchmarks with 8 threads, static schedule,  
passive wait-policy, simulated at 100 KIPS.

# Multi-threaded Sampling is Complex

Instruction count-based  
techniques are unsuitable<sup>1</sup>

Threads progress differently  
due to load imbalance

Representing parallelism  
among threads

Differentiating thread  
waiting from real work

# Multi-threaded Sampling is Complex

Instruction count-based  
techniques are unsuitable<sup>1</sup>

Threads progress differently  
due to load imbalance


Identify a *unit of work* that is *invariant* across executions

Representing parallelism  
among threads

Differentiating thread  
waiting from real work


# Extending Single-threaded Techniques

- SimPoint or SMARTS ➤ Instruction count-based techniques
  - Works well for single-threaded applications




# Extending Single-threaded Techniques

- SimPoint or SMARTS ➤ Instruction count-based techniques
  - Works well for single-threaded applications




# Extending Single-threaded Techniques

- SimPoint or SMARTS ➤ Instruction count-based techniques
  - Inconsistent regions for multi-threaded applications



# Extending Single-threaded Techniques


- SimPoint or SMARTS ➤ Instruction count-based techniques
  - Inconsistent regions for multi-threaded applications



# Multi-threaded Sampling

## FlexPoints


- + Designed for non-synchronizing throughput workloads
- ✓ Instruction count-based sampling
- ✗ Assumes no thread interaction
- ✗ Requires simulation of the full application



# Multi-threaded Sampling

## Time-based Sampling

- + Designed for synchronizing generic multi-threaded workloads
- ✓ Applies to generic multi-threaded workloads
- ✗ Extremely slow
- ✗ Requires simulation of the full application



# Multi-threaded Sampling

## BarrierPoint





Designed for barrier-synchronized multi-threaded workloads



Scales well with number of barriers



Slow when *inter-barrier regions* are large



# Multi-threaded Sampling

## TaskPoint



Designed for task-based workloads




Uses analytical models to improve accuracy



Works only for the particular workload type

```
#pragma omp task  
label(task_type 1)  
do_something();
```



# The Unit of Work

## Single-threaded Sampling

SimPoint<sup>1</sup>      ➡      Instruction count  
SMARTS<sup>2</sup>

## Multiprocessor Sampling

Flex Points<sup>3</sup>      ➡      Instruction count

## Multi-threaded Sampling

Time-based sampling<sup>4</sup>      ➡      Time

## Application-specific Sampling

BarrierPoint<sup>5</sup>      ➡      Inter-barrier regions  
TaskPoint<sup>6</sup>      ➡      Task instances

<sup>1</sup>Sherwood et al., “Automatically Characterizing Large Scale Program Behavior”, ASPLOS’02

<sup>2</sup>Wunderlich et al., “SMARTS: Accelerating Microarchitecture Simulation via Rigorous Statistical Sampling”, ISCA’03

<sup>3</sup>Wenisch et al., “SimFlex: statistical sampling of computer system simulation”, IEEE Micro’06

<sup>4</sup>Carlson et al., “Sampled Simulation of Multithreaded Applications”, ISPASS’13

<sup>5</sup>Carlson et al., “BarrierPoint: Sampled simulation of multi-threaded applications”, ISPASS’14

<sup>6</sup>Grass et al., “TaskPoint: Sampled simulation of task-based programs”, ISPASS’16

# The Unit of Work

## Single-threaded Sampling

SimPoint<sup>1</sup>  
SMARTS<sup>2</sup> → Instruction count

## Multiprocessor Sampling

Flex Points<sup>3</sup> → Instruction count

**We consider generic loop iterations as the unit of work**

Time-based sampling<sup>4</sup> → Time

BarrierPoint<sup>5</sup> → Inter-barrier regions  
TaskPoint<sup>6</sup> → Task instances

<sup>1</sup>Sherwood et al., "Automatically Characterizing Large Scale Program Behavior", ASPLOS'02

<sup>2</sup>Wunderlich et al., "SMARTS: Accelerating Microarchitecture Simulation via Rigorous Statistical Sampling", ISCA'03

<sup>3</sup>Wenisch et al., "SimFlex: statistical sampling of computer system simulation", IEEE Micro'06

<sup>4</sup>Carlson et al., "Sampled Simulation of Multithreaded Applications", ISPASS'13

<sup>5</sup>Carlson et al., "BarrierPoint: Sampled simulation of multi-threaded applications", ISPASS'14

<sup>6</sup>Grass et al., "TaskPoint: Sampled simulation of task-based programs", ISPASS'16

# Overall Methodology

Where to simulate


How to simulate

# Overall Methodology




How to simulate

# Overall Methodology




How to simulate


# Overall Methodology




# Overall Methodology





# Loop-based Profiling





# Loop-based Profiling



# Loop-based Profiling



# Loop-based Profiling



# Loop-based Profiling

Flow-control

Slice Generation  
(PC, count)

Synchronization  
Filtering

# Loop-based Profiling: Flow-control

- Load Imbalance can affect profiling
  - Make sure threads make equal forward progress
- Implementation: Control the forward progress of threads
  - Synchronize threads (barriers) externally at regular intervals
  - Make sure all threads execute similar number of instructions



Flow-control

Slice Generation  
(PC, count)

Synchronization  
Filtering



# Loop-based Profiling: Flow-control

- Load Imbalance can affect profiling
  - Make sure threads make equal forward progress
- Implementation: Control the forward progress of threads
  - Synchronize threads (barriers) externally at regular intervals
  - Make sure all threads execute similar number of instructions



# Loop-based Profiling: Flow-control

- Load Imbalance can affect profiling
  - Make sure threads make equal forward progress
- Implementation: Control the forward progress of threads
  - Synchronize threads (barriers) externally at regular intervals
  - Make sure all threads execute similar number of instructions





# Loop-based Profiling: Sync Filtering

- Goal: Filter out synchronization during profiling
  - Profiling data should contain only *real* work

- Solutions

- Automatic
- Ignore s



# Loop-based Profiling: Sync Filtering

- Goal: Filter out synchronization during profiling
  - Profiling data should contain only *real* work
- Solutions
  - Automatic detection using loop analysis<sup>1</sup>
  - Ignore sync library code (Ex. libiomp5.so, libpthread.so)


Flow-control

Slice Generation  
(PC, count)


Synchronization  
Filtering

# Loop-based Profiling: Sync Filtering

Ignore sync library code (Ex. libiomp5.so, libpthread.so)




Application execution




# Loop-based Profiling: Sync Filtering

Ignore sync library code (Ex. libiomp5.so, libpthread.so)




Application execution





# Loop-based Profiling: Slice Generation

- Region start/stop
  - Global instruction count reaches threshold ( $\#threads \times 100 \text{ M}$ )
  - Region boundary at a loop entry/exit – use DCFG analysis
- Looppoint region markers ( $PC, count_{PC}$ )
  - Global count of loop entries: invariant across executions
  - Simulate the same *amount of work*





# Loop-based Profiling: Slice Generation

- Region start/stop
  - Global instruction count reaches threshold ( $\#threads \times 100\text{ M}$ )
  - Region boundary at a loop entry/exit – use DCFG analysis
- Looppoint region markers ( $PC, count_{PC}$ )
  - Global count of loop entries: invariant across executions
  - Simulate the same *amount of work*



# Loop-based Profiling: Slice Generation

- Region start/stop
  - Global instruction count reaches threshold ( $\#threads \times 100 \text{ M}$ )
  - Region boundary at a loop entry/exit – use DCFG analysis
- Looppoint region markers ( $PC, count_{PC}$ )
  - Global count of loop entries: invariant across executions
  - Simulate the same *amount of work*



# Loop-based Profiling: Slice Generation

- Basic Block (BB)
  - Section of code with single entry and exit
- Basic Block Vector (BBV)
  - Execution fingerprint of an application interval
  - Vector with one element for each basic block
  - Exec Wt = *entry count* × *number of instructions*

ID:      A      B      C

BB	Example Assembly Code	
A	srl	a2, 0x8, t4
	and	a2, 0xff, t12
	addl	zero, t12, s6
	subl	t7, 0x1, t7
	cmpeq	s6, 0x25, v0
	cmpeq	s6, 0, t0
	bis	v0, t0, v0
	bne	v0, 0x120018c48
B	subl	t7, 0x1, t7
	cmple	t7, 0x3, t2
	beq	t2, 0x120018b04
C	ble	t7, 0x120018bb4
...	...	...

# Loop-based Profiling: Slice Generation

- Basic Block (BB)
  - Section of code with single entry and exit
- Basic Block Vector (BBV)
  - Execution fingerprint of an application interval
  - Vector with one element for each basic block
  - $\text{Exec Wt} = \text{entry count} \times \text{number of instructions}$

ID: A B C  
BB Exec Count: < 1, 20, 0, ...>  
weigh by Block Size: < 8, 3, 1, ...>

BB	Example Assembly Code
A	srl a2, 0x8, t4 and a2, 0xff, t12 addl zero, t12, s6 subl t7, 0x1, t7 cmpeq s6, 0x25, v0 cmpeq s6, 0, t0 bis v0, t0, v0 bne v0, 0x120018c48
B	subl t7, 0x1, t7 cmple t7, 0x3, t2 beq t2, 0x120018b04
C	ble t7, 0x120018bb4
...	...

# Loop-based Profiling: Slice Generation

- Basic Block (BB)
  - Section of code with single entry and exit
- Basic Block Vector (BBV)
  - Execution fingerprint of an application interval
  - Vector with one element for each basic block
  - Exec Wt = *entry count* × *number of instructions*

ID:    A    B    C

BB Exec Count: < 1, 20, 0, ...>

weigh by Block Size: < 8, 3, 1, ...>

BB Exec Wt: < 8, 60, 0, ...>

BB	Example Assembly Code
A	srl a2, 0x8, t4 and a2, 0xff, t12 addl zero, t12, s6 subl t7, 0x1, t7 cmpeq s6, 0x25, v0 cmpeq s6, 0, t0 bis v0, t0, v0 bne v0, 0x120018c48
B	subl t7, 0x1, t7 cmple t7, 0x3, t2 beq t2, 0x120018b04
C	ble t7, 0x120018bb4
...	...

# Loop-based Profiling: Slice Generation

- Basic Block (BB)
  - Section of code with single entry and exit
- Basic Block Vector (BBV)
  - Exec Count: < 1, 20, 0, ...>
  - Vector: [ A:8, B:60, C:0, ... ] BBV
  - Weigh by Block Size: < 8, 3, 1, ...>
  - Exec Wt = *entry count × number of instructions*

ID:    A    B    C

BB Exec Count: < 1, 20, 0, ...>




weigh by Block Size: < 8, 3, 1, ...>

BB Exec Wt: < 8, 60, 0, ...>

BB	Example Assembly Code
A	srl a2, 0x8, t4 and a2, 0xff, t12 addl zero, t12, s6 subl t7, 0x1, t7 cmpeq s6, 0x25, v0 cmpeq s6, 0, t0 bis v0, t0, v0 bne v0, 0x120018c48
B	subl t7, 0x1, t7 cmple t7, 0x3, t2 beq t2, 0x120018b04
C	ble t7, 0x120018bb4
...	...

# Loop-based Profiling: Vector Concatenation

- Ratio of instructions per thread may differ
- *Global-BBVs*: Concatenate per-thread BBVs to larger Global BBV



# Loop-based Profiling: Vector Concatenation

- Ratio of instructions per thread may differ
- Global-BBVs:** Concatenate per-thread BBVs to larger Global BBV

BB	Example Assembly Code		
A	srl	a2, 0x8, t4	
	and	a2, 0xff, t12	
	addl	zero, t12, s6	
	subl	t7, 0x1, t7	
	cmpeq	s6, 0x25, v0	
	cmpeq	s6, 0, t0	
	bis	v0, t0, v0	
	bne	v0, 0x120018c48	
B	subl	t7, 0x1, t7	
	cmple	t7, 0x3, t2	
	beq	t2, 0x120018b04	
C	ble	t7, 0x120018bb4	
...	...		
M	subl	t7, 0x1, t7	
	gt	t7, 0x120018b90	

BB	Example Assembly Code		
A	srl	a2, 0x8, t4	
	and	a2, 0xff, t12	
	addl	zero, t12, s6	
	subl	t7, 0x1, t7	
	cmpeq	s6, 0x25, v0	
	cmpeq	s6, 0, t0	
	bis	v0, t0, v0	
	bne	v0, 0x120018c48	
B	subl	t7, 0x1, t7	
	cmple	t7, 0x3, t2	
	beq	t2, 0x120018b04	
C	ble	t7, 0x120018bb4	
...	...		
M	subl	t7, 0x1, t7	
	gt	t7, 0x120018b90	

# Loop-based Profiling: Vector Concatenation

- Ratio of instructions per thread may differ
- Global-BBVs*: Concatenate per-thread BBs into global BBV

BB	Example Assembly Code			
A				
	BB	Example Assembly Code		
	A	srl	a2, 0x8, t4	
		and	a2, 0xff, t12	
		addl	zero, t12, s6	
		subl	t7, 0x1, t7	
		cmpeq	s6, 0x25, v0	
		cmpeq	s6, 0, t0	
		bis	v0, t0, v0	
		bne	v0, 0x120018c48	
B				
	B	subl	t7, 0x1, t7	
		cmple	t7, 0x3, t2	
		beq	t2, 0x120018b04	
C				
...	C	ble	t7, 0x120018bb4	
M		...		
	M	subl	t7, 0x1, t7	
		gt	t7, 0x120018b90	

# Loop-based Profiling: Vector Concatenation

- Ratio of instructions per thread may differ
- Global-BBVs*: Concatenate per-thread BBs into global BBV

BB ID:      A      B      C      ...  
BB Exec Wt: < 8,    60,    0, ... >

BB ID:      N      O      P      ...  
BB Exec Wt: < 5,    90,    3, ... >

BB	Example Assembly Code			
N				
A	srl a2, 0x8, t4 and a2, 0xff, t12 addl zero, t12, s6 subl t7, 0x1, t7 cmpeq s6, 0x25, v0 cmpeq s6, 0, t0 bis v0, t0, v0 bne v0, 0x120018c48			
O				
B	subl t7, 0x1, t7 cmple t7, 0x3, t2 beq t2, 0x120018b04			
P				
...	ble t7, 0x120018bb4			
Z	...			
M	subl t7, 0x1, t7 gt t7, 0x120018b90			

Thread 1

Thread 0

# Loop-based Profiling: Vector Concatenation

- Ratio of instructions per thread may differ
- Global-BBVs*: Concatenate per-thread BBs into Global-BBV

BB ID:      A      B      C      ...  
BB Exec Wt: < 8,    60,    0,    ... >

The diagram illustrates the concatenation of per-thread basic blocks (BBs) into a global basic block (BBV). It shows two tables of assembly code, one for Thread 1 (blue dashed box) and one for Thread 0 (red dashed box). The tables have columns for BB ID and Example Assembly Code.

**Thread 1 (Blue Dashed Box):**

BB	Example Assembly Code
N	srl a2, 0x8, t4 and a2, 0xff, t12 addl zero, t12, s6 subl t7, 0x1, t7 0x25, v0 0, t0 t0, v0 0x120018c48

**Thread 0 (Red Dashed Box):**


BB	Example Assembly Code
B	subl t7, 0x1, t7 cmple t7, 0x3, t2 beq t2, 0x120018b04
C	ble t7, 0x120018bb4
P	...
Z	...
M	subl t7, 0x1, t7 gt t7, 0x120018b90

A blue callout box highlights the "Global-BBV" row, which contains the concatenated assembly code for both threads: [ A:8, B:60, C:0, ..., N:5, O:90, P:3, ... ].

BB ID:      N      O      P      ...  
BB Exec Wt: < 5,    90,    3,    ... >

# A LoopPoint Region


```
638.imagick_s/magick/morphology.c
2842 #if defined(MAGICKCORE_OPENMP_SUPPORT)
2843     #pragma omp parallel for schedule(static,4) shared(progress,status) \
2844         magick_threads(image,result_image,image->rows,1)
2845 #endif
2846     for (y=0; y < (ssize_t) image->rows; y++)
2847     {
2848         ...
2849         for (x=0; x < (ssize_t) image->columns; x++)
2850         {
2851             for (v=0; v < (ssize_t) kernel->height; v++) {
2852                 for (u=0; u < (ssize_t) kernel->width; u++, k--) {
2853                     ...
2854                     } /* u */
2855                     ...
2856                     } /* v */
2857                 } /* x */
2858             } /* y */
2859             ...
2860     }
```



638.imagick\_s, train input, 8 threads

# A LoopPoint Region


```
638.imagick_s/magick/morphology.c
2842 #if defined(MAGICKCORE_OPENMP_SUPPORT)
2843     #pragma omp parallel for schedule(static,4) shared(progress,status) \
2844         magick_threads(image,result_image,image->rows,1)
2845 #endif
2846     for (y=0; y < (ssize_t) image->rows; y++)
2847     {
2848         ...
2849         for (x=0; x < (ssize_t) image->columns; x++)
2850         {
2851             ...
2852             for (v=0; v < (ssize_t) kernel->height; v++) {
2853                 for (u=0; u < (ssize_t) kernel->width; u++, k--) {
2854                     ...
2855                     } /* u */
2856                     ...
2857                     } /* v */
2858                 } /* x */
2859             } /* y */
2860             ...
```



638.imagick\_s, train input, 8 threads


# Identifying Simulation Regions

- Group similar Global-BBVs
  - K-means algorithm: Centroid-based clustering
- Vector closest to centroid is the representative
- Simulation regions (looppoints)
  - Checkpoints generated from the application
  - Use  $(PC, \text{count}_{PC})$  information of representatives




# Identifying Simulation Regions

- Group similar Global-BBVs
  - K-means algorithm: Centroid-based clustering
- Vector closest to centroid is the representative
- Simulation regions (looppoints)
  - Checkpoints generated from the application
  - Use  $(PC, \text{count}_{PC})$  information of representatives




# Identifying Simulation Regions

- Group similar Global-BBVs
  - K-means algorithm: Centroid-based clustering
- Vector closest to centroid is the representative
- Simulation regions (looppoints)
  - Checkpoints generated from the application
  - Use  $(PC, \text{count}_{PC})$  information of representatives



# Identifying Simulation Regions

- Group similar Global-BBVs
  - K-means algorithm: Centroid-based clustering
- Vector closest to centroid is the representative
- Simulation regions (looppoints)
  - Checkpoints generated from the application
  - Use  $(PC, \text{count}_{PC})$  information of representatives



# Application Reconstruction

- Representative regions (looppoints) are simulated in parallel
- Warmup handling
  - Simulate a large enough warmup region before simulation region
- Application performance
  - The weighted average of the performance of simulation regions



# Application Reconstruction

- Representative regions (looppoints) are simulated in parallel
- Warmup handling
  - Simulate a large enough warmup region before simulation region
- Application performance
  - The weighted average of the performance of simulation regions



$$\text{total runtime} = \sum_{i=rep_1}^{rep_N} \text{runtime}_i \times \text{multiplier}_i$$

# Application Reconstruction

- Representative regions (looppoints) are simulated in parallel
- Warmup handling
  - Simulate a large enough warmup region before simulation region
- Application performance
  - The weighted average of the performance




$$multiplier_j = \frac{\sum_{i=0}^m inscount_i}{inscount_j}$$

$m$  regions represented by  $j^{\text{th}}$  looppoint

$$\text{total runtime} = \sum_{i=rep_1}^{rep_N} \text{runtime}_i \times \text{multiplier}_i$$

# Experimental Setup

- Simulation Infrastructure
  - Sniper<sup>1</sup> 7.4
    - Mimics Intel Gainestown 8/16 core
- Benchmarks and OpenMP settings
  - SPEC CPU2017 speed benchmarks
    - Input: train; Threads: 8; Wait policy: Active, Passive
  - NAS Parallel Benchmarks (NPB)
    - Input: Class C; Threads: 8, 16; Wait policy: Passive
  - OpenMP scheduling policy: *static*



# SPEC CPU2017 Analysis


Application (speed version)	Parallel	static for	dyna mic for	barrier (explic it)	master	single	reduction (nowait)	atomic (float8_a dd)	atomic (float8 _max)	atomic (fixed4_ add)	lock
603.bwaves	Yes	Yes					Yes	Yes	Yes		
607.cactuBSSN	Yes	Yes	Yes	Yes			Yes	Yes			
619.lbm	Yes	Yes									
621.wrf	Yes		Yes		Yes						
627.cam4	Yes	Yes	Yes	Yes	Yes						
628.pop2	Yes	Yes		Yes	Yes						
638.imagick	Yes	Yes		Yes	Yes	Yes					Yes
644.nab	Yes		Yes	Yes			Yes	Yes		Yes	
649.fotonik3d	Yes	Yes									
654.roms	Yes	Yes									

# Workload Type Supported

- Software
  - Static OpenMP scheduling (`OMP_WAIT_POLICY=STATIC`)
  - Homogeneous parallel threads doing equal amount of work
- Hardware
  - Simulated hardware needs to be homogeneous
  - No dynamic hardware events supported


# Accuracy Results

Prediction error wrt. performance of whole application




# Accuracy Results

Prediction error wrt. performance of whole application



# Accuracy Results


Prediction error wrt. performance of whole application



# Changing Thread Count

Runtime prediction error wrt. whole application runtime


NPB 3.3 with *Class C* inputs, 8 and 16 threads, *passive* wait-policy



# Changing Thread Count

Runtime prediction error wrt. whole application runtime


NPB 3.3 with *Class C* inputs, 8 and 16 threads, *passive wait-policy*



# Speedup

Parallel and serial speedup achieved for LoopPoint


SPEC CPU2017 with *train* inputs, 8 threads, *active* wait-policy



# Speedup

Parallel and serial speedup achieved for LoopPoint


SPEC CPU2017 with *train* inputs, 8 threads, *active* wait-policy




# Speedup

## Parallel and serial speedup achieved for LoopPoint

SPEC CPU2017 with *train* inputs, 8 threads, active wait-policy




NPB with *Class C* inputs, 8 and 16 threads, passive wait-policy




# Speedup

## Parallel and serial speedup achieved for LoopPoint

SPEC CPU2017 with *train* inputs, 8 threads, *active* wait-policy




NPB with *Class C* inputs, 8 and 16 threads, *passive* wait-policy



# Speedup

## Theoretical Speedup comparison with BarrierPoint


SPEC CPU2017 with *ref* inputs, 8 threads, *passive* wait-policy



# Speedup

## Theoretical Speedup comparison with BarrierPoint

SPEC CPU2017 with *ref* inputs, 8 threads, *passive* wait-policy



Up to 31000X  
speedup!

Serial: 244x  
Parallel: 11587x

# Summary

- Contributions
  - Methodology to sample generic multi-threaded workloads
  - Uses application loops (barring spinloops) as the unit of work
  - Flexible to be used for checkpoint-based simulation
- Accurate results in minimal time
  - Average absolute error of 2.3% across applications
  - Parallel speedup going up to 31,000 ×
  - Reduces simulation time from a few years to a few hours

# More Information

- Links
  - Artifact: <https://github.com/nus-comparch/looppoint>
  - Page: <https://looppoint.github.io>
  - Short talk: <https://youtu.be/Tr609MkT42g>
  - Questions: [alens@comp.nus.edu.sg](mailto:alens@comp.nus.edu.sg), [tcarlson@comp.nus.edu.sg](mailto:tcarlson@comp.nus.edu.sg)
- Upcoming tutorial session ➤ *LoopPoint and ELFies*
  - ISCA 2022, New York City



# Agenda

Time	Speaker	Topic
09.00 to 09.10	Alen Sabu	Overview of the tutorial
09.10 to 10.30	Harish Patil	Tools from Intel: Pin, PinPlay, SDE, ELFies
10.30 to 10.45	Break	
10.45 to 11.30	Akanksha Chaudhari	Simulation and Single-threaded Sampling
11.30 to 11.40	Break	
11.40 to 12.20	Alen Sabu	Multi-threaded Sampling and LoopPoint
12.20 to 13.00	Changxi Liu	Running Sniper and LoopPoint Tools

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation




Session 4

# Sniper and LoopPoint Demo

CHANGXI LIU, PHD STUDENT

NATIONAL UNIVERSITY OF SINGAPORE

# Simulator Design Waterfall




- Cycle-accurate simulation is too slow
- High-level simulation consider both accuracy and execution time

# Sniper: A Fast and Accurate Simulator


- Hybrid simulation approach
  - Analytical interval core model : Interval Model
  - Micro-architecture structure simulation
    - Branch predictors, caches, etc.
- Support multi/many-cores simulation with parallel scales of core number
- Pin-based frontend, can also support dynamoRIO
- Open source <https://snipersim.org>

# Interval Model



- Split whole application into consecutive intervals

# Simulation in Sniper



# Directly Running Sniper

- Download
  - <https://snipersim.org/w/Download>
  - Register to receive the link via email.

- External Email -

Dear Sniper downloader,

Here are your download instructions for the Sniper Multi-core Simulator.

For use with GIT, clone our repository from the following directory:

```
$ git clone http://snipersim.org/download/ [REDACTED]/git/sniper.git
```

To download Sniper, use the following link or command:

```
http://snipersim.org/download/ [REDACTED]/packages/sniper-latest.tgz  
$ wget http://snipersim.org/download/ [REDACTED]/packages/sniper-latest.tgz
```

If you have any questions, feel free to post them on our mailing list

<http://groups.google.com/group/snipersim>

or visit our Frequently Asked Questions page

[http://snipersim.org/w/Frequently\\_Asked\\_Questions](http://snipersim.org/w/Frequently_Asked_Questions)

The Sniper Simulator Team

# Directly Running Sniper

- Download
  - <https://snipersim.org/w/Download>
  - Register to receive the link via email.
- Simulating an application
  - `./run-sniper <options> -- /bin/ls`

- External Email -

Dear Sniper downloader,

Here are your download instructions for the Sniper Multi-core Simulator.

For use with GIT, clone our repository from the following directory:

`$ git clone http://snipersim.org/download/[REDACTED]/git/sniper.git`

To download Sniper, use the following link or command:

[http://snipersim.org/download/\[REDACTED\]/packages/sniper-latest.tgz](http://snipersim.org/download/[REDACTED]/packages/sniper-latest.tgz)  
`$ wget http://snipersim.org/download/[REDACTED]/packages/sniper-latest.tgz`

If you have any questions, feel free to post them on our mailing list

<http://groups.google.com/group/snipersim>

or visit our Frequently Asked Questions page

[http://snipersim.org/w/Frequently\\_Asked\\_Questions](http://snipersim.org/w/Frequently_Asked_Questions)

The Sniper Simulator Team

# Build LoopPoint

- Prerequisites
  - X86-based Linux machine
  - C++ with c++11 support
  - Python
  - Docker
  - Sniper link

# Build LoopPoint

- Opensource code
  - <https://github.com/nus-comparch/looppoint.git>

# Build LoopPoint

- make build
  - Build docker image

```
or Python 2.7 in January 2021. More details about Python 2 support in pip can be found at https://pip.pypa.io/en/latest/development/release-process/#python-2-support
p 21.0 will remove support for this functionality.
Collecting tabulate
  Downloading tabulate-0.8.9.tar.gz (53 kB)
Building wheels for collected packages: tabulate
  Building wheel for tabulate (setup.py): started
  Building wheel for tabulate (setup.py): finished with status 'done'
    Created wheel for tabulate: filename=tabulate-0.8.9-py2-none-any.whl size=33171 sha256=c5713bdcee7e07c602619a643e2c3132e8e25d18308dc1d0fe06a9ef8b03e12
    Stored in directory: /tmp/pip-ephem-wheel-cache-N3SSxV/wheels/0a/4b/e1/d0e504a346ed0882b93f971fe1122b9de64fabebd9b1d81b9f
Successfully built tabulate
Installing collected packages: tabulate
Successfully installed tabulate-0.8.9
Removing intermediate container 94d7acee39fa
--> 98fe9327b5dc
Step 9/9 : RUN pip3 install --no-cache-dir --upgrade pip &&     pip3 install --no-cache-dir numpy
--> Running in 3b699c6d695a
Collecting pip
  Downloading https://files.pythonhosted.org/packages/a4/6d/6463d49a933f547439d6b5b98b46af8742cc03ae83543e4d7688c2420f8b/pip-21.3.1-py3-none-any.whl (1.7MB)
Installing collected packages: pip
  Found existing installation: pip 9.0.1
    Not uninstalling pip at /usr/lib/python3/dist-packages, outside environment /usr
Successfully installed pip-21.3.1
WARNING: pip is being invoked by an old script wrapper. This will fail in a future version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip directly.
Collecting numpy
  Downloading numpy-1.19.5-cp36-cp36m-manylinux2010_x86_64.whl (14.8 MB)
Installing collected packages: numpy
Successfully installed numpy-1.19.5
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended to use a virtual environment instead: https://pip.pypa.io/warnings/venv
Removing intermediate container 3b699c6d695a
--> 57f0a752e1e6
[Warning] One or more build-args [TZ_ARG] were not consumed
Successfully built 57f0a752e1e6
Successfully tagged ubuntu:18.04-looppoint
```

# Build LoopPoint

- make build
- make
  - Run the docker image

```
I have no name!@ef5546e12134:~/spass/looppoint$ ls  
Dockerfile-ubuntu-18.04 README.md lplib.py run-looppoint.py tools  
Makefile apps preprocess suites.py  
I have no name!@ef5546e12134:~/spass/looppoint$
```

# Build LoopPoint

- make build
- make
- make apps
  - Build the provided application
  - matrix-mul demo
  - You can find the source code of the demo in
    - apps/demo/matrix-omp/
  - **Coming soon:** Support for open-source benchmarks (like NPB) with LoopPoint

```
I have no name!@ef5546e12134:~/ispass/looppoint$ make apps
make -C apps/demo/matrix-omp
make[1]: Entering directory '/home/ispass/looppoint/apps/demo/matrix-omp'
g++ -g -O3 -fopenmp -o matrix-omp matrix-omp-init.cpp matrix-omp.cpp -static
/usr/lib/gcc/x86_64-linux-gnu/7/libgomp.a(target.o): In function `gomp_target'
(.text+0x8b): warning: Using 'dlopen' in statically linked applications requ
ln -s matrix-omp base.exe
make[1]: Leaving directory '/home/ispass/looppoint/apps/demo/matrix-omp'
I have no name!@ef5546e12134:~/ispass/looppoint$ ls
Dockerfile-ubuntu-18.04 Makefile README.md apps lplib.py preprocess ru
I have no name!@ef5546e12134:~/ispass/looppoint$ find -name base.exe
./apps/demo/matrix-omp/base.exe
I have no name!@ef5546e12134:~/ispass/looppoint$
```

# Build LoopPoint

- make build
- make
- make apps
- make tools SNIPER\_GIT\_REPO=[1]
  - Build Sniper and LoopPoint tools

```
I have no name!@ef5546e12134: [REDACTED] spass/looppoint$ make tools SNIPER_GIT_REPO=http://snipersim.org/download/[REDACTED]/git/sniper.git
Downloading Pin
--2022-05-17 07:34:17-- https://software.intel.com/sites/landingpage/pintool/downloads/pin-3.13-98189-g60a6ef199-gcc-linux.tar.gz
Resolving software.intel.com (software.intel.com)... 23.15.96.37, 2600:1417:3f:791::b, 2600:1417:3f:78d::b
Connecting to software.intel.com (software.intel.com)|23.15.96.37|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 35897895 (34M) [application/octet-stream]
Saving to: 'STDOUT'

[REDACTED] 100%[=====] 34.23M 11.4MB/s in 3.0s
```

```
2022-05-17 07:34:20 (11.4 MB/s) - written to stdout [35897895/35897895]

patching file pin-3.13-98189-g60a6ef199-gcc-linux/source/tools/InstLib/alarms.H
patching file pin-3.13-98189-g60a6ef199-gcc-linux/source/tools/InstLib/alarms.cpp
Setting SNIPER_GIT_REPO as http://snipersim.org/download/[REDACTED]/git/sniper.git to download Sniper
Cloning into 'tools/sniper'...
```

```
[CXX ] sift/recorder/obj-intel64/thread.o
[CXX ] sift/recorder/obj-intel64/papi.o
[CXX ] sift/recorder/obj-intel64/bbv_count.o
[CXX ] sift/recorder/obj-intel64/trace rtn.o
[CXX ] sift/recorder/obj-intel64/recorder_base.o
[CXX ] sift/recorder/obj-intel64/pinboost_debug.o
[CXX ] sift/recorder/obj-intel64/syscall_modeling.o
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder/sift'
[CXX ] sift/recorder/sift/sift_reader.o
[CXX ] sift/recorder/sift/zfstream.o
[CXX ] sift/recorder/sift/sift_utils.o
[CXX ] sift/recorder/sift/sift_writer.o
[LD ] sift/recorder/sift/libshift.a
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder/sift'
[LD ] sift/recorder/obj-intel64/sift_recorder
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/sift/recorder'
make[4]: Entering directory '[REDACTED] /spass/looppoint/tools/sniper/standalone'
make[4]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/standalone'
[DEP ] standalone/standalone.d
[DEP ] standalone/exceptions.d
[CXX ] standalone/exceptions.o
[CXX ] standalone/standalone.o
[LD ] lib/sniper
make[2]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper/standalone'
make[1]: Leaving directory '[REDACTED] /spass/looppoint/tools/sniper'
I have no name!@ef5546e12134: [REDACTED] spass/looppoint$
```

# Build LoopPoint

- Opensource code
- We provide the script to help you build the environment
  - make build
    - Build docker image
  - make
    - Run docker image
  - make apps
    - Build the provided applications
  - make tools SNIPER\_GIT\_REPO=[1]
    - Build Sniper and LoopPoint tools

# Running LoopPoint

- Then run LoopPoint!
  - `./run-looppoint.py -h`
  - Provides the information on how to run the tool
- Example run command
  - `./run-looppoint.py -p demo-matrix-1 -n 8 --force`

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
    - make\_mt\_pinball : Generate whole-program pinball
    - gen\_dcfg : Generate DCFG file to identify loop information
    - gen\_bbv : Generate feature vector of each region
    - gen\_cluster : Cluster regions

# PinPlay

- Makes Pin-based analyses repeatable.
- Command:
  - `$SDE_KIT/pinplay-scripts/sde_pinpoints.py --mode mt --cfg=$CFGFILE --log_options="-start_address main -log:fat -log:basename $WPP_BASE" --replay_options="-replay:strace" -l`
- Generates a whole-program pinball for further profiling steps

# DCFG

- A control-flow graph (CFG) is a fundamental structure
- A dynamic control-flow graph (DCFG) is a specialized CFG that adds data from a specific execution of a program
- C++ DCFG APIs is conveniently to used for accessing the data.
  - `DCFG_LOOP_CONTAINER::get_loop_ids`
    - Get the set of loop IDs
  - `DCFG_LOOP`
    - `get_routine_id` : get the function that the loop belongs to
    - `get_parent_loop_id` : get the parent loop

# DCFG

- A control-flow graph (CFG) is a fundamental structure
- A dynamic control-flow graph (DCFG) is a specialized CFG that adds data from a specific execution of a program
- C++ DCFG APIs is conveniently to used for accessing the data.
- More APIs can be found in
  - `tools/sde-external-9.0.0-2021-11-07-lin/pinkit/sde-example/include`
    - `dcfg_api.H`
    - `dcfg_pin_api.H`
    - `dcfg_trace_api.H`

# DCFG

- Collect Loop Information
- Command:
  - `$SDE_BUILD_KIT/pinplay-scripts/replay.py --pintool=sde-global-looppoint.so --pintool_options "-dcfg" -replay:deadlock_timeout 0 -replay:strace -dcfg:out_base_name $DCFG_BASE $WPP_BASE"`
  - -dcfg : enable DCFG generation
  - DCFG\_BASE : the DCFG file name that is generated

- Profiling the feature vector of each region
- Command:
  - `$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pintool="sde-global-looppoint.so" --global_regions --pccount_regions --cfg $CFG --whole_pgm_dir $WPP_DIR --mode mt -S $SLICESIZE -b --replay_options "-replay:deadlock_timeout 0 -global_profile -emit_vectors 0 -filter_exclude_lib libgomp.so.1 -filter_exclude_lib libiomp5.so -looppoint:global_profile -looppoint:dcfg-file $DCFG -looppoint:main_image_only 1 -looppoint:loop_info $PROGRAM.$INPUT.loop_info.txt -flowcontrol:verbose 1 -flowcontrol:quantum 1000000 -flowcontrol:maxthreads $NCORES"`
  - `--pccount_regions` : (PC, count)-based region information
  - `--looppoint:loop_info` : Utilize loop information as the marker of each region
  - `--flowcontrol:quantum` : synchronize each thread every 1000000 instructions

# Clustering

- Cluster all regions into several groups.
  - SimPoint [1]
  - Utilize feature vectors of all threads
  - kmeans algorithm

# Clustering

- Cluster all regions into several groups.
- Command
  - `$SDE_BUILD_KIT/pinplay-scripts/sde_pinpoints.py --pintool="sde-global-looppoint.so" --cfg $CFG --whole_pgm_dir $WPP_DIR -S $SLICESIZE --warmup_factor=2 --maxK=$MAXK --append_status -s --simpoint_options="-dim $DIM -coveragePct 1.0 -maxK $MAXK"`
  - DIM : The reduced dimension of the vector that BBVs are projected to
  - MAXK : Maximum number of clusters for kmeans

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling Final Results:
    - matrix.1\_16448.global.pinpoints.csv
    - (start-pc, start-pc-count ), (end-pc, end-pc-count )

```
^ 3 # comment,thread-id,region-id,start-pc, start-image-name, start-image-offset, start-pc-count,end-pc, end-image-name, end-image-offset, end-pc-count, region-length, region-weight, region-multiplier, region-type
4
5 # RegionId = 1 Slice = 0 Icount = 0 Length = 800000067 Weight = 0.12500 Multiplier = 1.000 ClusterSliceCount = 1 ClusterIcount = 800000066
6 #Start: pc : 0x400880 image: matrix-omp offset: 0x880 absolute_count: 1 source-info: matrix-omp.cpp:17
7 #End: pc : 0x4040c0 image: matrix-omp offset: 0x40c0 absolute_count: 77977888 relative_count: 9837476.0 source-info: matrix-omp.cpp:75
8 cluster 0 from slice 0,global,1,0x400880,matrix-omp,0x880,1,0x4040c0,matrix-omp,0x40c0,77977888,9837476,800000067,0.12500,1.000,simulation
```

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling Final Results:
    - matrix.1\_16448.global.pinpoints.csv
    - (start-pc, start-pc-count ), (end-pc, end-pc-count)
    - Cluster group id

```
^ 3 # comment,thread-id,region-id,start-pc, start-image-name, start-image-offset, start-pc-count,end-pc, end-image-name, end-image-offset, end-p  
count, region-length, region-weight, region-multiplier, region-type  
4  
5 # RegionId = 1 Slice = 0 Icount = 0 Length = 800000067 Weight = 0.12500 Multiplier = 1.000 ClusterSliceCount = 1 ClusterIcount = 800000066  
6 #Start: pc : 0x400880 image: matrix-omp offset: 0x880 absolute_count: 1 source-info: matrix-omp.cpp:17  
7 #End: pc : 0x4040c0 image: matrix-omp offset: 0x40c0 absolute_count: 77977888 relative_count: 9837476.0 source-info: matrix-omp.cpp:75  
8 cluster 0 from slice 0,global,1,0x400880,matrix-omp,0x880,1,0x4040c0,matrix-omp,0x40c0,77977888,9837476,800000067,0.12500,1.000,simulation  
^
```

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling Final Results:
    - matrix.1\_16448.global.pinpoints.csv
    - (start-pc, start-pc-count ), (end-pc, end-pc-count)
    - Cluster group id
    - Cluster multiplier

```
^ 3 # comment,thread-id,region-id,start-pc, start-image-name, start-image-offset, start-pc-count,end-pc, end-image-name, end-image-offset, end-p
   count, region-length, region-weight, region-multiplier, region-type
 4
 5 # RegionId = 1 Slice = 0 Icount = 0 Length = 800000067 Weight = 0.12500 Multiplier = 1.000 ClusterSliceCount = 1 ClusterIcount = 800000066
 6 #Start: pc : 0x400880 image: matrix-omp offset: 0x880 absolute_count: 1 source-info: matrix-omp.cpp:17
 7 #End: pc : 0x4040c0 image: matrix-omp offset: 0x40c0 absolute_count: 77977888 relative_count: 9837476.0 source-info: matrix-omp.cpp:75
 8 cluster 0 from slice 0,global,1,0x400880,matrix-omp,0x880,1,0x4040c0,matrix-omp,0x40c0,77977888,9837476,800000067,0.12500 [1.000] simulation
```

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
    - `matrix.1_16448.global.pinpoints.csv`
    - Sampled Simulation : (start-pc, start-pc-count ), (end-pc, end-pc-count), cluster group id
    - Extrapolation : cluster group id, cluster-multiplier

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
  - Sampled simulation of selected regions

# Sniper

- The LoopPoint support in Sniper
  - Handle the beginning and ending of representative regions

```
VOID Handler(CONTROLLER::EVENT_TYPE ev, VOID * v, CONTEXT * ctxt
             VOID * ip, THREADID tid, BOOL bcast)
{
    switch(ev)
    {
        case CONTROLLER::EVENT_START:
            handleMagic(tid, ctxt, SIM_CMD_USER, 0x0be0000f, 0);
            break;

        case CONTROLLER::EVENT_STOP:
            handleMagic(tid, ctxt, SIM_CMD_USER, 0x0be0000f, 1);
            break;

        default:
            break;
    }
}
```

# Sniper

- The LoopPoint support in Sniper
  - Handle the beginning and ending of representative regions
  - Register this function in pin

```
VOID Handler(CONTROLLER::EVENT_TYPE ev, VOID * v, CONTEXT * ctxt
             VOID * ip, THREADID tid, BOOL bcast)
{
    switch(ev)
    {
        case CONTROLLER::EVENT_START:
            handleMagic(tid, ctxt, SIM_CMD_USER, 0x0be0000f, 0);
            break;

        case CONTROLLER::EVENT_STOP:
            handleMagic(tid, ctxt, SIM_CMD_USER, 0x0be0000f, 1);
            break;

        default:
            break;
    }
}

control_manager.RegisterHandler(Handler, 0, FALSE);
control_manager.Activate();
```

# Sniper

- The LoopPoint support in Sniper
  - Handle the beginning and ending of representative regions
  - Register this function in pin
  - ```
./run-sniper -n 8 -gscheduler/type=static -cgainestown -ssimuserroi --roi-script --trace-args=-control start:address:0x4069d0:count235036646:global --trace-args=-control stop:address:0x4069d0:count313177121:global -- <app cmd>
```

    - -control start:address:<PC>:<Count>
    - -control end:address:<PC>:<Count>
  - PC , Count :LoopPoint region boundaries

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
  - Sampled simulation of selected regions
  - Extrapolation of performance results

# Extrapolation of Performance Result

- Runtime of corresponding representative region : **regionid**
- Multiply the ratio : **multiplier**

```
for regionid, multiplier in region_mult.iteritems():
    region_runtime = 0
    try:
        region_runtime = read_simstats(region_stats[regionid], region_config[regionid], 'runtime')
    except:
        print('[LOOPPOINT] Warning: Skipping r%s as the simulation results are not available' % regionid)
        continue
    cov_mult += multiplier
    extrapolated_runtime += region_runtime * multiplier
    if region_runtime > max_rep_runtime:
        max_rep_runtime = region_runtime
        sum_rep_runtime += region_runtime
```

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
  - Sampled simulation of selected regions
  - Extrapolation of performance results
    - Predicted runtime using sampled simulation

| application  | runtime<br>actual (ns) | runtime<br>predicted (ns) | error<br>(%) | speedup<br>(parallel) | speedup<br>(serial) | coverage<br>(%) |
|--------------|------------------------|---------------------------|--------------|-----------------------|---------------------|-----------------|
| matrix-omp.1 | 214544900.0            | 199674000.0               | 6.93         | 8.34                  | 4.24                | 100.0           |

# Running LoopPoint

- The driver script of LoopPoint
  - Profiling the application
  - Sampled simulation of selected regions
  - Extrapolation of performance results
    - Predicted runtime using sampled simulation
    - The error rate of obtained using sampled simulation

| application  | runtime<br>actual (ns) | runtime<br>predicted (ns) | error<br>(%) | speedup<br>(parallel) | speedup<br>(serial) | coverage<br>(%) |
|--------------|------------------------|---------------------------|--------------|-----------------------|---------------------|-----------------|
| matrix-omp.1 | 214544900.0            | 199674000.0               | 6.93         | 8.34                  | 4.24                | 100.0           |

Thank you!

# LoopPoint and ELFies: Tools and Techniques to Accelerate Simulations of Multi-threaded Applications using Checkpointing

Alen Sabu<sup>1</sup>, Changxi Liu<sup>1</sup>, Akanksha Chaudhari<sup>1</sup>, Harish Patil<sup>2</sup>,  
Wim Heirman<sup>2</sup>, Trevor E. Carlson<sup>1</sup>

<sup>1</sup>National University of Singapore

<sup>2</sup>Intel Corporation

